

■ Cuttlefish Memory Doctrine & MCP Integration

The Cuttlefish Labs AI architecture embraces Model Context Protocol (MCP) as a core infrastructure for maintaining

"Insanity is a breakdown in persistent context—so is misaligned AI."

By using MCP, each agent gains a cryptographically verifiable thread of reasoning, allowing for trusted memory

Key Features:

1. MCPClientContext in every CuttlefishAgent to track tool usage and decision history.
2. Long-Term Memory Registry embedded in Cortex layer to manage goal states, failures, and hallucination risk.
3. Reflection loops governed by `.ask()` / `.remember()` / `.verify()` for ethical alignment.
4. Secure AuthKit token flow to ensure valid agent action with scope-bound access.

This integration unlocks:

- Agent identity continuity
- Multi-agent planning
- Investor-grade trust and rollback

This is the first cognitive firewall for sovereign infrastructure intelligence.